

CS536

Introduction

to

Advanced Compiler

A Sahu
CSE, IIT Guwahati

<http://jatinga.iitg.ac.in/~asahu/cs536/>

Outline

- Course Structure
- Pre-requisite
- Reference Books
- Grading Policy and Assignments
- Introduction and Motivation to course

<http://jatinga.iitg.ac.in/~asahu/cs536/>

Why to study compiler?

Impact!

Techniques in compilers helps all programmers

Compiler technology

- **Bridge gap between HLL and Machine**

Concept of Prog. Lang.

- High Level Prog. Lang.
- Domain Specific Lang.
- Natural Lang.

Concept of Comp. Arch.

- RISC vs CISC
- Locality: Cache, memory hierarchy
- Parallelism : Instruction level, Thread level

Programmers



Prog. Language

Compilers



Machine
Code

Prog. Tools

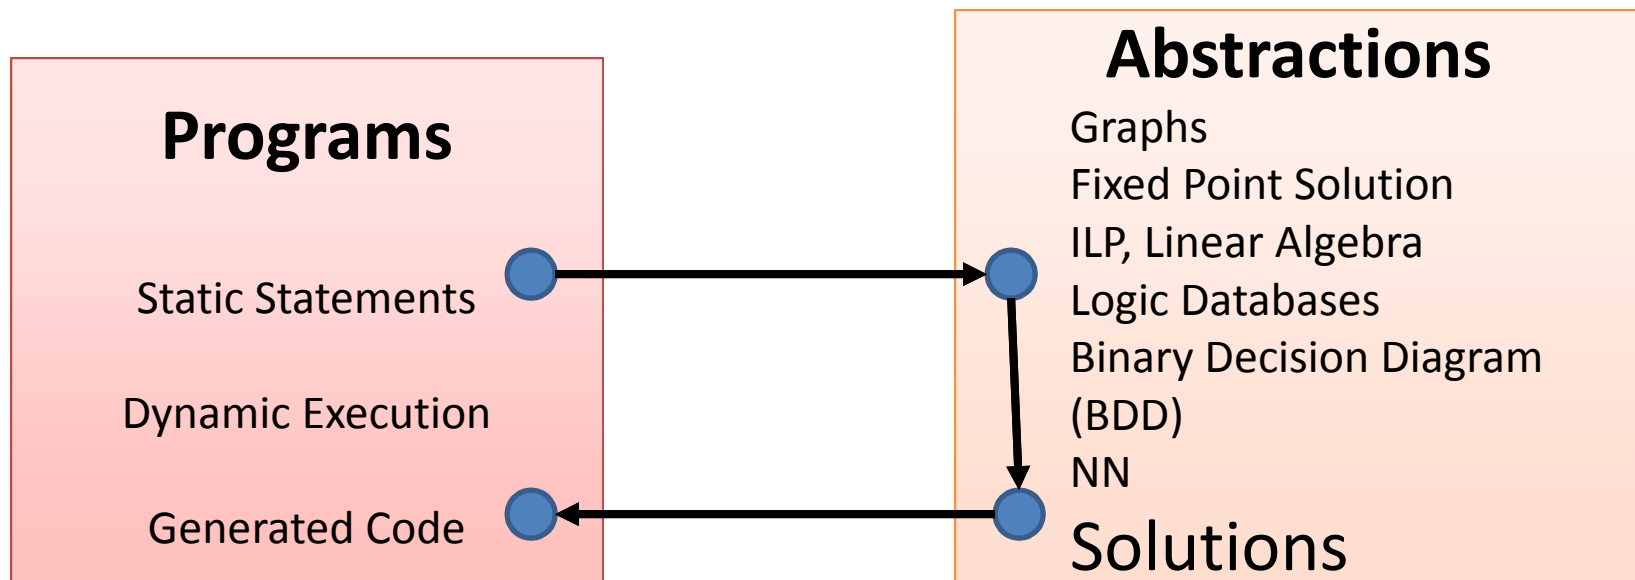
- Binary Translation
- Security

Compiler Study Trains Good Developers

- Reasoning about programs makes better programmers
- Tool building: there are programmers and there are tool builders
- Excellent software engineering case study: Compilers are hard to build
 - Input: all programs
 - Objectives:
- Methodology for solving complex real -life problems
 - Build upon mathematical / programming abstractions

Compilers: Where theory meets practice

- Desired solutions are often NP-complete /undecidable
- Key to success: Formulate right abstraction/ approx.
 - Can't be solved by just pure hacking
 - theory aids generality and correctness
 - Can't be solved by just theory :
 - Expt. validates & provides feedback to problem formulation
- Tradeoffs: Generality, power, simplicity, and efficiency



Compiler Study Trains Good Developers

- Reasoning about programs makes better programmers
- Tool building: there are programmers and there are tool builders
- Excellent software engineering case study: Compilers are hard to build
 - Input: all programs
 - Objectives:
- Methodology for solving complex real -life problems
 - Build upon mathematical / programming abstractions


How Can the Compiler Improve Performance?

Execution time = Operation count * Cycles per operation

- Minimize the number of operations
 - arithmetic operations, memory accesses
- Replace expensive operations with simpler ones
 - e.g., replace 4 - cycle multiplication with 1 - cycle shift
- Minimize cache misses
 - both data and instruction accesses
- Perform work in parallel
 - instruction scheduling within a thread
 - parallel execution across multiple threads

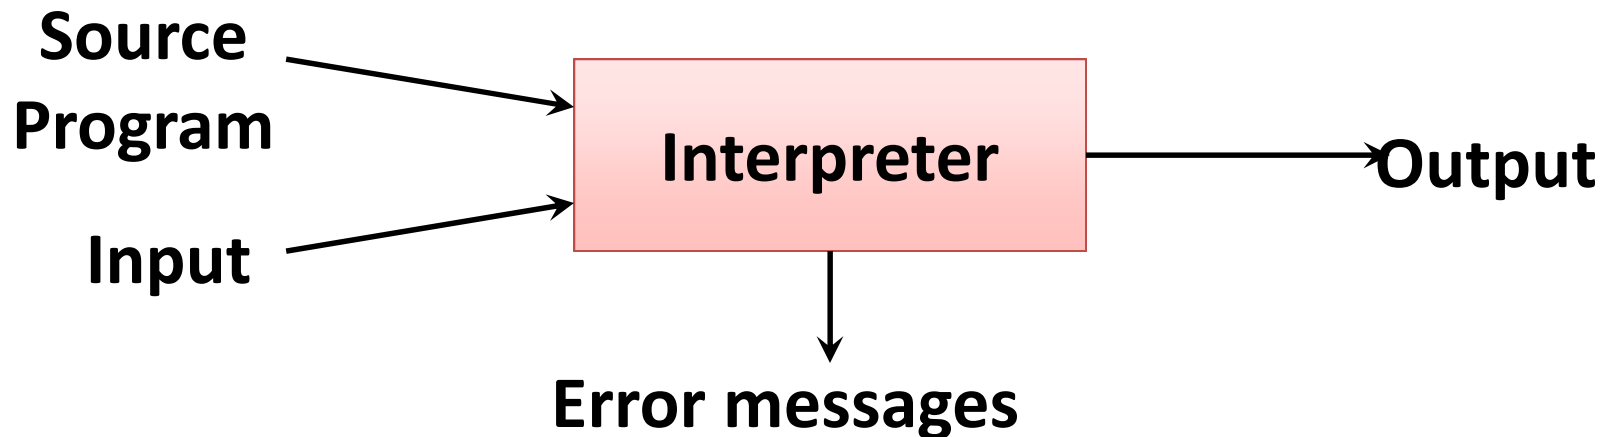
Compilers and Interpreters

- “*Compilation*”
 - Translation of a program written in a source language into a semantically equivalent program written in a target language



Compilers and Interpreters (cont'd)

- “*Interpretation*”
 - Performing the operations implied by the source program



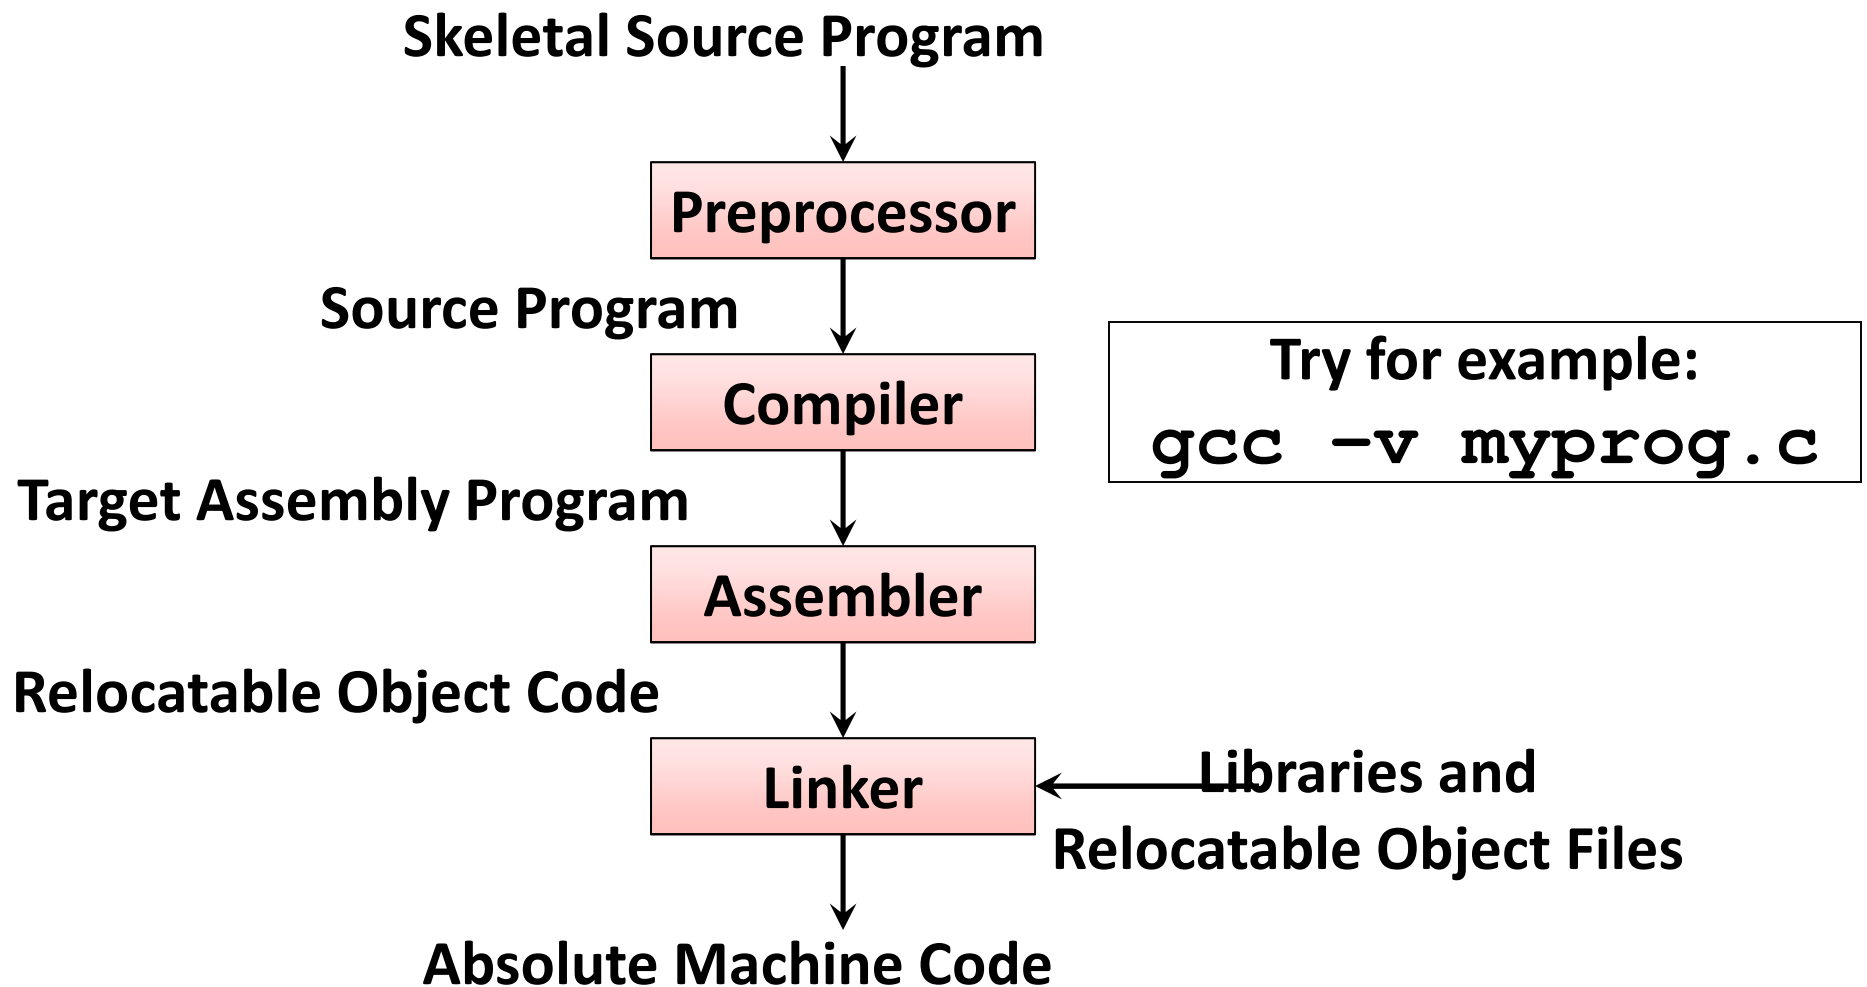
The Analysis-Synthesis Model of Compilation

- There are two parts to compilation:
 - *Analysis* determines the operations implied by the source program which are recorded in a tree structure
 - *Synthesis* takes the tree structure and translates the operations therein into the target program

Other Tools that Use the Analysis-Synthesis Model

- Editors (syntax highlighting)
- Pretty printers (e.g. doxygen)
- Static checkers (e.g. lint and splint)
- Interpreters
- Text formatters (e.g. TeX and LaTeX)
- Silicon compilers (e.g. VHDL)
- Query interpreters/compilers (Databases)

Preprocessors, Compilers, Assemblers, and Linkers



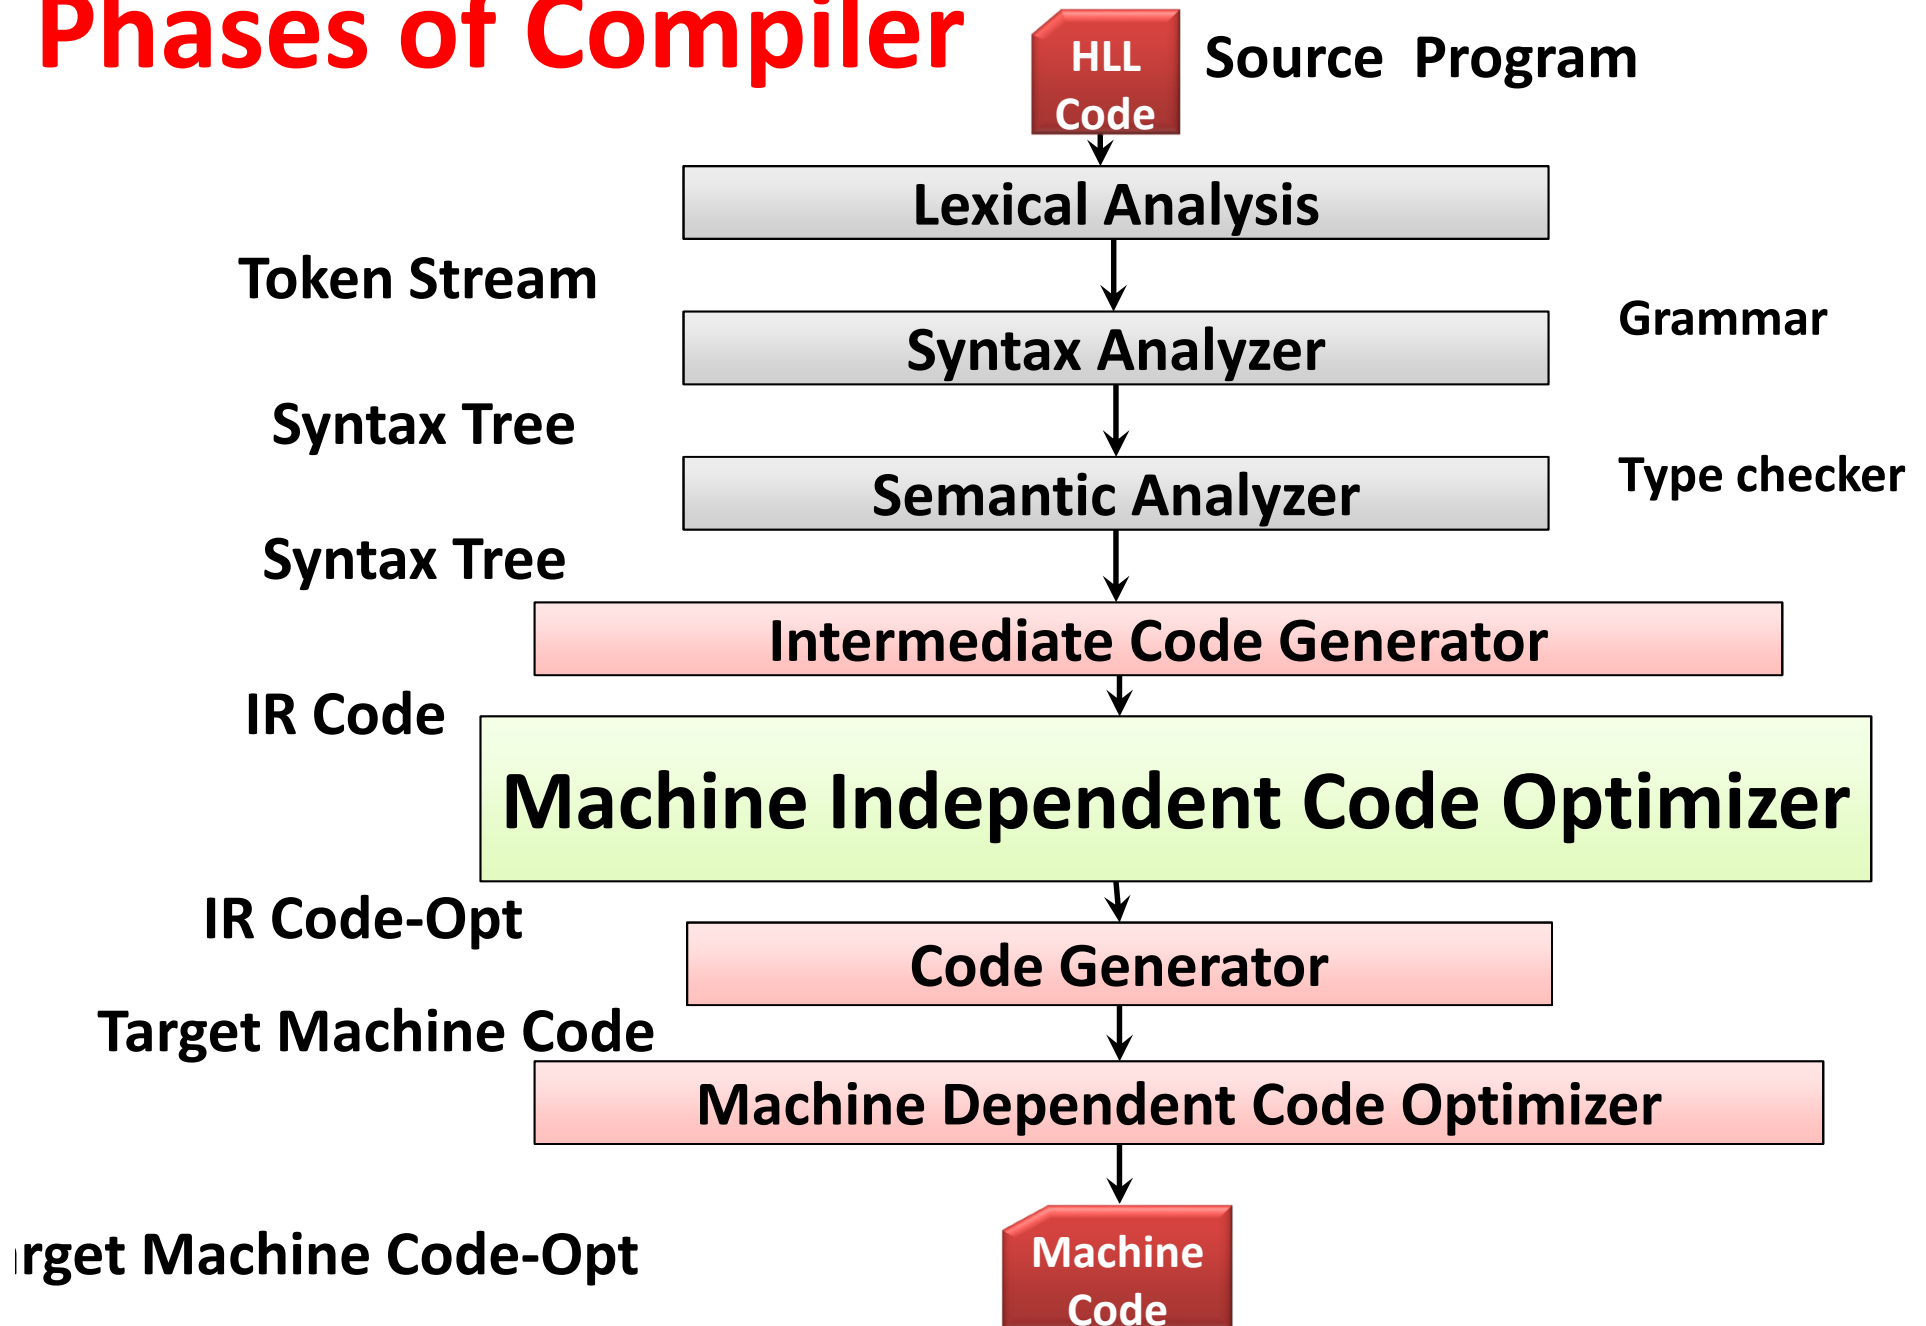
RISC vs CISC

- RISC : arithmetic and memory instruction are well separated
 - RISC : All instruction have same size
 - RISC: Less number of Addressing mode
 - RISC : Less number of type of instruction
- ➔
- Simple/less power/less area controller
 - Pipeline able design for RISC
 - Example of RISC: ARM (M1 from Apple, SD, MT), MIPS, PowePC
- *Example of CISC : X86/x86-64, VxWorks*
 - *CISC use micro-code to convert complex INS to simple INS*

GCC Options

- `$gcc -E test.c // Pre-process file`
- `$gcc -save-temps t.c`
`//save Preprocess file (t.i), t.s (assembly), t.o`

Phases of Compiler



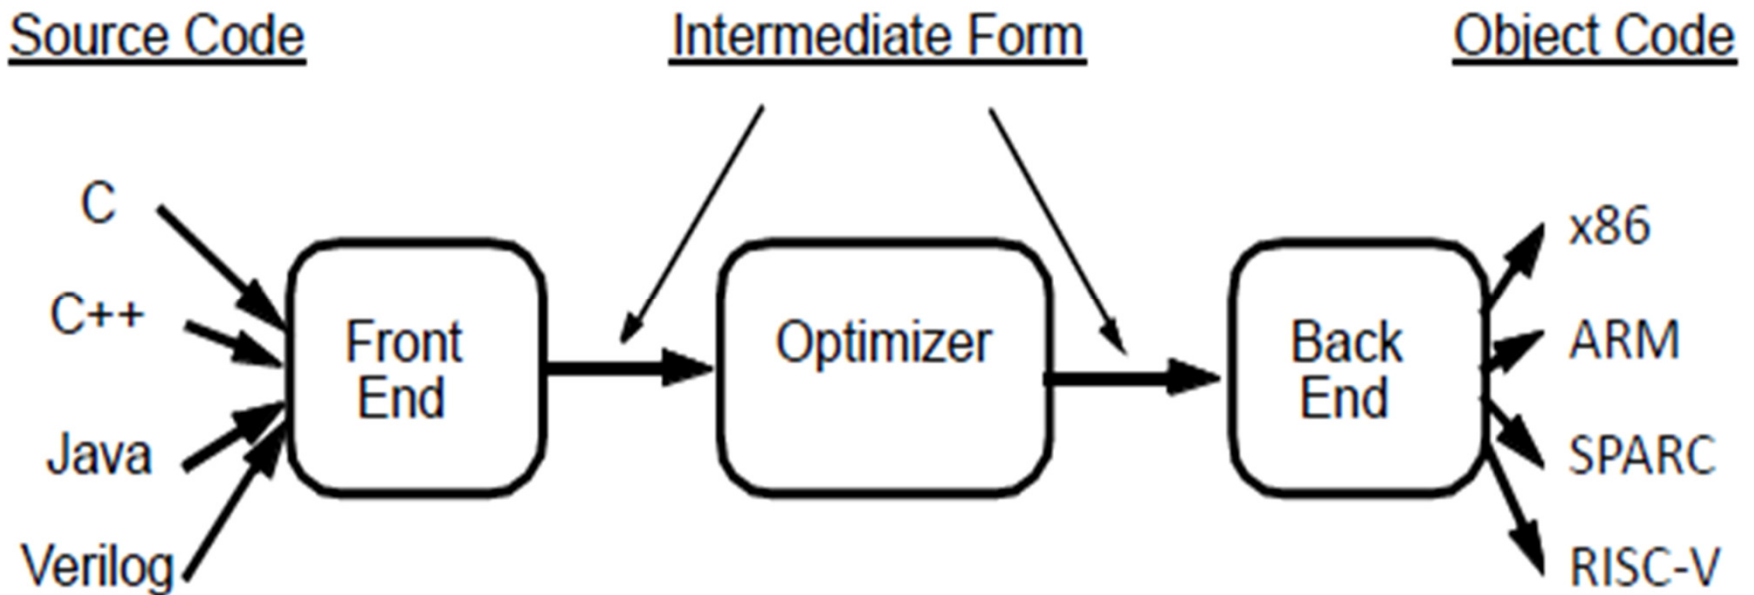
The Phases of a Compiler

Phase	Output	Sample
<i>Programmer</i>	Source string	A=B+C;
<i>Scanner</i> (performs <i>lexical analysis</i>)	Token string	'A', '=', 'B', '+', 'C', ';'' And <i>symbol table</i> for identifiers
<i>Parser</i> (performs <i>syntax analysis</i> based on the grammar of the programming language)	Parse tree or abstract syntax tree	<pre> ; = / \ A + / \ B C </pre>
<i>Semantic analyzer</i> (type checking, etc)	Parse tree or abstract syntax tree	
<i>Intermediate code generator</i>	Three-address code, quads, or RTL	<pre> int2fp B t1 + t1 C t2 := t2 A </pre>
<i>Optimizer</i>	Three-address code, quads, or RTL	<pre> int2fp B t1 + t1 #2.3 A </pre>
<i>Code generator</i>	Assembly code	<pre> MOVF #2.3, r1 ADDF2 r1, r2 MOVF r2, A </pre>
<i>Peephole optimizer</i>	Assembly code	<pre> ADDF2 #2.3, r2 MOVF r2, A </pre>

Grouping of Phases

- Compiler front and back ends:
 - Analysis (*machine independent* front end)
 - Synthesis (*machine dependent* back end)
- Passes
 - A collection of phases may be repeated only once (*single pass*) or multiple times (*multi pass*)
 - Single pass: usually requires everything to be defined before being used in source program
 - Multi pass: compiler may have to keep entire program representation in memory

Structure of a Compiler



- Optimizations are performed on an “intermediate form ”
 - similar to a generic RISC instruction set
- Enables easy portability to multiple source languages, target machines

General GCC Optimization Options

- GCC optimization : - O0, - O1, -O2, -O3
- `$man gcc`
- At – O0 level:
 - Compiler refrain from most of the opt.
 - It is correct choice for analyzing the code with debugger
- At high level
 - Mixed up source lines, eliminate redundant variable, rearrange arithmetic expressions
 - Debugger has a hard time to give user a consistent view on code and data

General GCC Optimization Options

- Level 1
 - -fauto-inc-dec, -fmove-loop-invariant, -fmerge-constants, -ftree-copy-prop, -finline-fun-called-once
- Level 2
 - -falign-functions, -falign-loops, level, -finlining-small-fun, -finling-indirect-fun, -freorder-fun, -fstrict-aliasing
- Level 3
 - -ftree-slp-vectorize, -fvect-cost-model

CS536 Course Flow

- Intermediate Representation
- IR Code Generation
- Runtime Environment
- Code Generation: Basic Block, SSA, Reg. Alloc.
- **Machine In/Dependent Code Optimization**
- **Optimization for ILP**
- **Opt. for Parallelism (TLP) and Cache Locality**
- Inter procedural Analysis

Motivation 1: Optimizing Compilers for High Level Prog. Language

Example: Bubblesort

program that sorts array A allocated in static storage

```
for (i = n - 2; i >= 0; I -- ) {  
  for (j = 0; j <= I ; j++) {  
    if (A[j] > A[j+1]) {  
      temp = A[j];  
      A[j] = A[j+1];  
      A[j+1] = temp;  
    }  
  }  
}
```

Code generated by front end

```

i := n-2
s5: if i<0 goto s1
j := 0
s4: if j>i goto s2
t1 = 4*j
t2 = &A
t3 = t2+t1
t4 = *t3          ;A[j]
t5 = j+1
t6 = 4*t5
t7 = &A
t8 = t7+t6
t9 = *t8          ;A[j+1]
if t4 <= t9 goto s3
t10 = 4*j
t11 = &A
t12 = t11+t10
temp = *t12      ;temp=A[j]

t13 = j+1
t14 = 4*t13
t15 = &A
t16 = t15+t14
t17 = *t16        ;A[j+1]
t18 = 4*j
t19 = &A
t20 = t19+t18     ;&A[j]
*t20 = t17        ;A[j]=A[j+1]
t21 = j+1
t22 = 4*t21
t23 = &A
t24 = t23+t22
*t24 = temp       ;A[j+1]=temp
s3: j = j+1
    goto s4
s2: i = i-1
    goto s5
s1:
```

*(t4=*t3 means read memory at address in t3 and write to t4,
t20=t17 means store value of t17 into memory at address in t20)

After optimization

Result of applying:

global common subexpression

loop invariant code motion

induction variable elimination

dead-code elimination


to all the scalar and temp. variables

These traditional optimizations can make
a big difference!

```
i = n-2
t27 = 4*i
t28 = &A
t29 = t27+t28
t30 = t28+4
s5: if t29 < t28 goto s1
t25 = t28
t26 = t30
s4: if t25 > t29 goto s2
t4 = *t25           ;A[j]
t9 = *t26           ;A[j+1]
if t4 <= t9 goto s3
temp = *t25         ;temp=A[j]
t17 = *t26          ;A[j+1]
*t25 = t17          ;A[j]=A[j+1]
*t26 = temp         ;A[j+1]=temp
s3: t25 = t25+4
t26 = t26+4
goto s4
s2: t29 = t29-4
goto s5
s1:
```

Motiv2: High Perf. Machine Learning


7 ExaFLOPS
60 Million Parameters



2015

Microsoft ResNet
Superhuman Image
Recognition

20 ExaFLOPS
300 Million Parameters



2016

Baidu Deep Speech 2
Superhuman Voice
Recognition

100 ExaFLOPS
8700 Million Parameters



2017

Google Neural
Machine Translation
Near Human
Language Translation

1 ExaFLOPS = 10^{18} FLOPS

2  NVIDIA

Motiv2:Nvidia Hopper GPU Arch



80B transistors
900 mm²
1650 MHz

132 Streaming Multiprocessors (SM)
528 Tensor cores
16,896 CUDA cores

Each SM



128 FP32 cores

128 Int cores

64 FP64 cores

4 Tensor cores

Tensor Cores

$D = A \times B + C$; A, B, C, D are 4x4 matrices

4 x 4 x 4 matrix processing array

1024 floating point ops / clock



FP64: 30 TFLOPS

FP32: 500 TFLOPS

FP16: 1000 TFLOPS


FP8/INT8 Tensor: 2000 TFLOPS

Blocking for Matrix Multiplication



Data
Accessed

1002000



65024

Blocking for Matmul: Original Code

```
for (i= 0; i< n; i++) {  
    for (j = 0; j < n; j++) {  
        for (k = 0; k < n; k++) {  
            Z[i,j] = Z[i,j] + X[i,k]*Y[k,j];  
        }  
    }  
}
```


Blocking for Matmul: Stripmine 2 outerloop

```
for (ii = 0; ii < n; ii = ii+B) {  
    for (i= ii; i< min(n,ii+B); i++) {  
        for (jj= 0; jj< n; jj= jj+B) {  
            for (j = jj; j < min(n,jj+B); j++) {  
                for (k = 0; k < n; k++) {  
                    Z[i,j] = Z[i,j] + X[i,k]*Y[k,j];  
                }  
            }  
        }  
    }  
}
```

Blocking for Matmul: permute

```
for (ii = 0; ii < n; ii = ii+B) {  
  for (jj= 0; jj< n; jj= jj+B) {  
    for (k = 0; k < n; k++) {  
      for (i= ii; i< min(n,ii+B); i++) {  
        for (j = jj; j < min(n,jj+B); j++) {  
          Z[i,j] = Z[i,j] + X[i,k]*Y[k,j];  
        }  
      }  
    }  
  }  
}
```


Blocking for Matmul :Impact



Thanks